

# Structured Streaming

# Learning Objectives

- ▶ Process streaming data
- ▶ `DataStreamReader`
- ▶ `DataStreamWriter`

# Data Stream

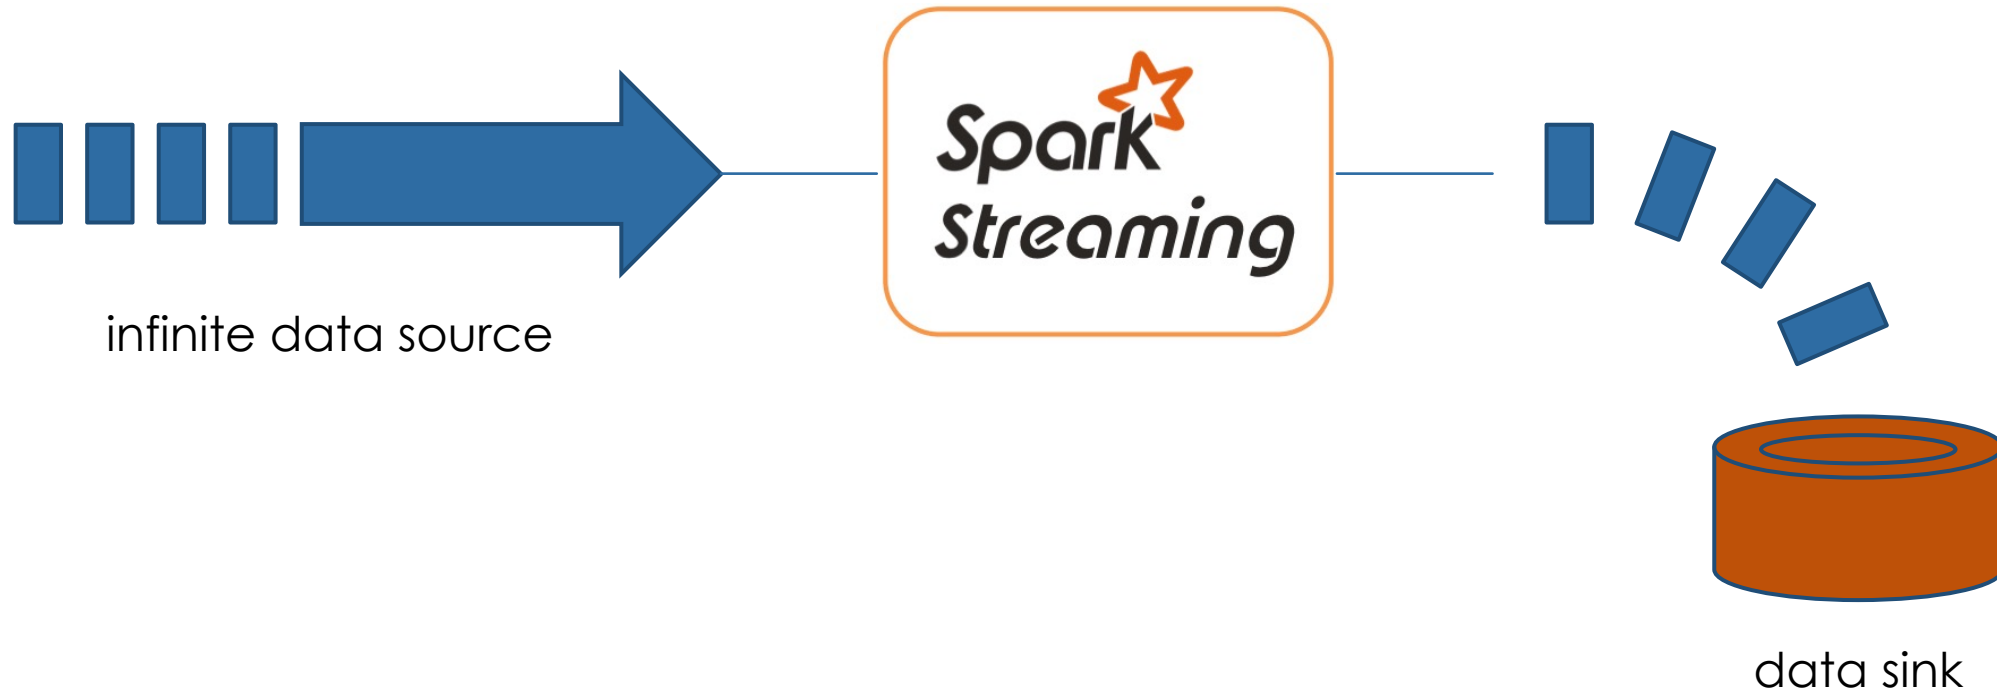
- ▶ Any data source that grows over time
- ▶ New files landing in cloud storage
- ▶ Updates to a database captured in a CDC feed
- ▶ Events queued in a pub/sub messaging feed

# Processing Data Stream

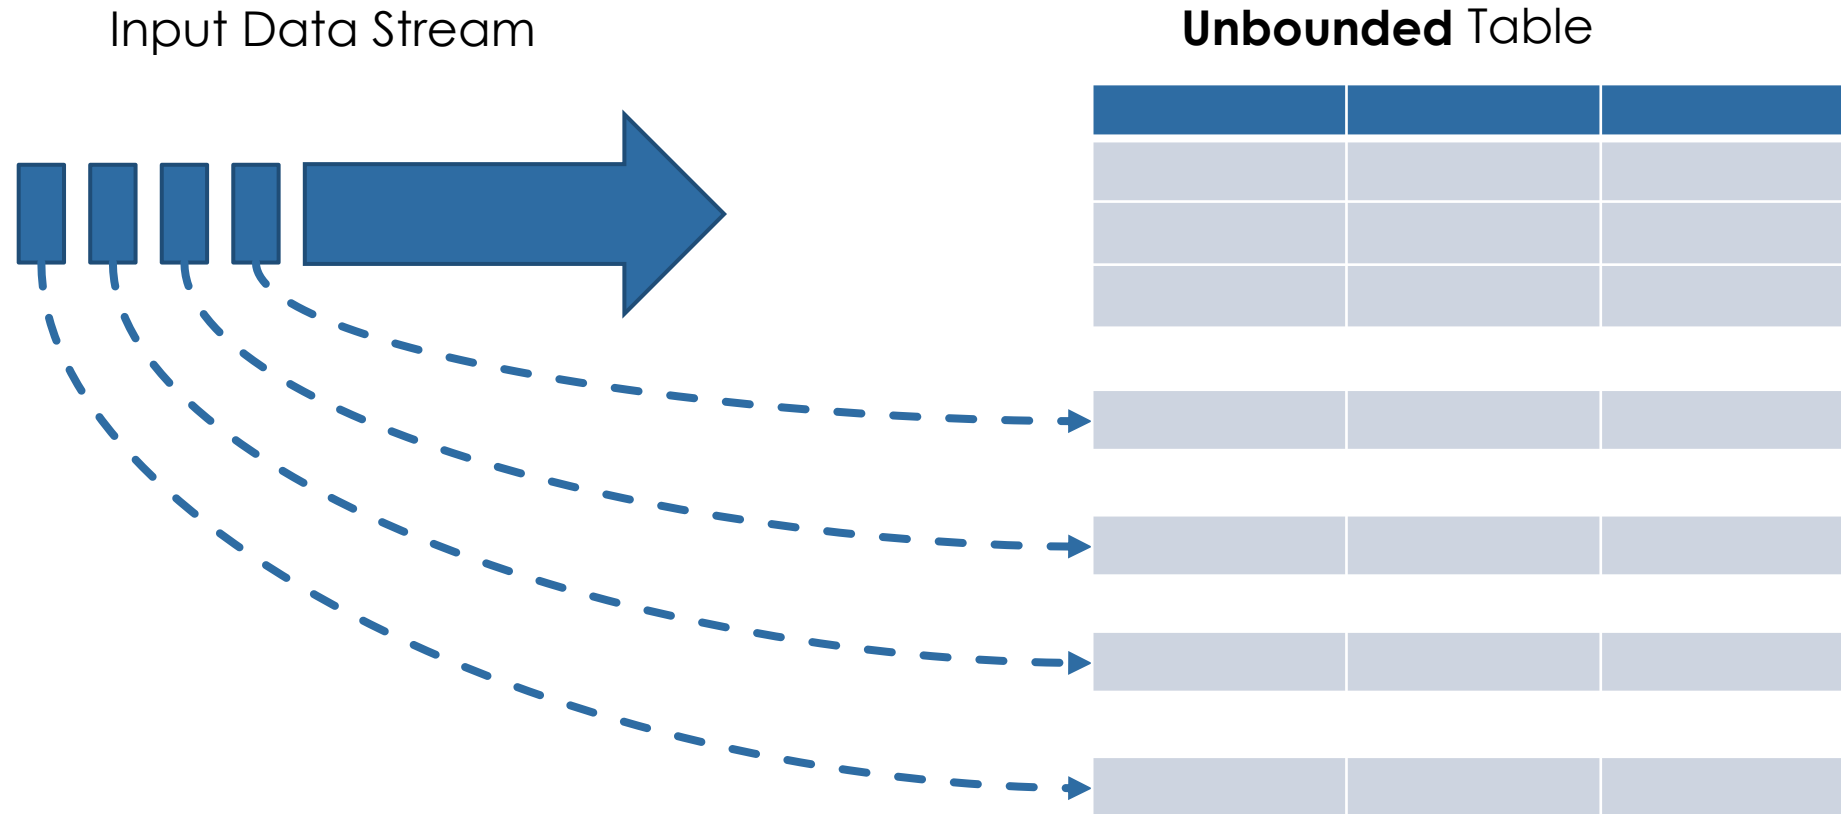
▶ 2 approaches:

1. Reprocess the entire source dataset each time
2. Only process those new data added since last update
  - ▶ Structured Streaming

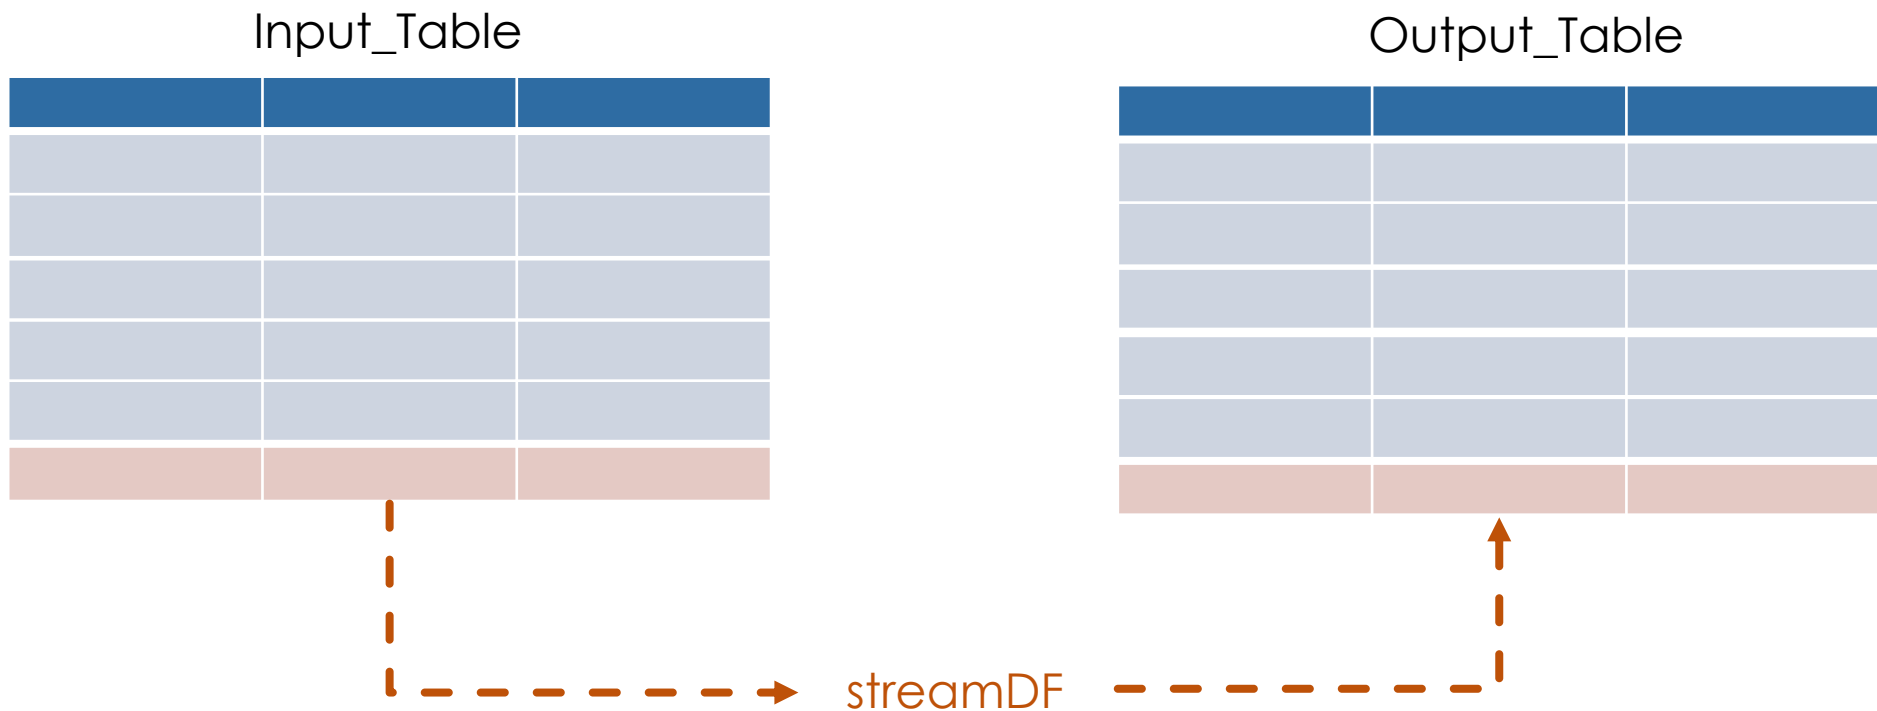
# Spark Structured Streaming



# Treating Infinite Data as a Table



# Input Streaming Table



```
streamDF = spark.readStream  
    .table("Input_Table")
```

```
streamDF.writeStream  
    .trigger(processingTime="2 minutes")  
    .outputMode("append")  
    .option("checkpointLocation", "/path")  
    .table("Output_Table")
```

# Trigger Intervals

```
streamDF.writeStream  
  .trigger(processingTime="2 minutes")  
  .outputMode("append")  
  .option("checkpointLocation", "/path")  
  .table("Output_Table")
```

Trigger	Method call	Behavior
Unspecified		Default: processingTime="500ms"
Fixed interval	.trigger( <b>processingTime</b> ="5 minutes")	Process data in micro-batches at the user-specified intervals
Triggered batch	.trigger( <b>once</b> =True)	Process all available data in a single batch, then stop
Triggered micro-batches	.trigger( <b>availableNow</b> =True)	Process all available data in multiple micro-batches, then stop



# Output Modes

```
streamDF.writeStream  
    .trigger(processingTime="2 minutes")  
    .outputMode("append")  
    .option("checkpointLocation", "/path")  
    .table("Output_Table")
```

Mode	Method call	Behavior
Append (Default)	<code>.outputMode("append")</code>	Only newly appended rows are incrementally appended to the target table with each batch
Complete	<code>.outputMode("complete")</code>	The target table is overwritten with each batch

# Checkpointing

```
streamDF.writeStream  
    .trigger(processingTime="2 minutes")  
    .outputMode("append")  
    .option("checkpointLocation", "/path")  
    .table("Output_Table")
```

- ▶ Store stream state
- ▶ Track the progress of your stream processing
- ▶ Can **Not** be shared between separate streams

# Guarantees

## 1. Fault Tolerance

- ▶ Checkpointing + Write-ahead logs
  - ▶ record the offset range of data being processed during each trigger interval.

## 2. Exactly-once guarantee

- ▶ Idempotent sinks

# Unsupported Operations

- ▶ Some operations are not supported by streaming DataFrame
  - ▶ Sorting
  - ▶ Deduplication
- ▶ Advanced methods
  - ▶ Windowing
  - ▶ Watermarking